

SQL Server Data File Architecture

Component	Description	Details
Data Files (.mdf, .ndf)	Physical files that store the data and objects of a SQL Server database.	<ul style="list-style-type: none"> - Primary Data File (.mdf): The main data file containing startup information for the database. - Secondary Data Files (.ndf): Optional files to spread data across multiple disks.
Filegroups	Logical grouping of data files within a database, allowing easier management and optimization.	<ul style="list-style-type: none"> - Primary Filegroup: Contains the primary data file (.mdf) and objects not assigned to other filegroups. - Secondary Filegroups: Used for partitioning or to distribute large tables and indexes across multiple files.
Extents	A collection of eight 8 KB pages, totaling 64 KB, which SQL Server uses to manage space allocation in data files.	<ul style="list-style-type: none"> - Uniform Extents: All eight pages belong to a single object. - Mixed Extents: Pages can belong to different objects, used for small tables or indexes.
Pages	The fundamental unit of data storage in SQL Server, each 8 KB in size.	<ul style="list-style-type: none"> - Types: Data Pages (storing table data), Index Pages, GAM/SGAM Pages (managing extents), PFS Pages (tracking page usage). - Page Header: Contains metadata like page ID, type, and status.
GAM and SGAM Pages	Specialized pages that manage and track the allocation of extents in the data files.	<ul style="list-style-type: none"> - GAM (Global Allocation Map): Tracks which extents are allocated. - SGAM (Shared Global Allocation Map): Tracks mixed extents that have free pages.
PFS Pages (Page Free Space)	Pages that record the allocation status and free space availability of individual pages within an extent.	<ul style="list-style-type: none"> - Purpose: Used to quickly find free space within an extent. - Frequency: A PFS page occurs every 8,088 pages (approximately 64 MB).
IAM Pages (Index Allocation Map)	Pages that map the extents used by a table or index, helping SQL Server manage space allocation for objects.	<ul style="list-style-type: none"> - Purpose: Tracks which extents belong to a table or index. - Multiple IAM Pages: Required for large tables or indexes spanning multiple filegroups or extents.

LOB Data Pages	Pages used to store large object (LOB) data types such as <code>text</code> , <code>ntext</code> , <code>image</code> , <code>varchar(max)</code> , <code>nvarchar(max)</code> .	<ul style="list-style-type: none"> - Storage: Stored separately from regular data pages if the data exceeds 8 KB. - Pointers: Regular pages store pointers to LOB data pages.
Row-Overflow Data	Handles variable-length columns that exceed the 8 KB page size limit.	<ul style="list-style-type: none"> - Overflow Pages: Data is moved to these pages when it exceeds the 8 KB limit. - Pointer: The original page stores a pointer to the row-overflow page.
Partitions	Logical divisions of a table or index across multiple filegroups, improving manageability and performance.	<ul style="list-style-type: none"> - Purpose: Helps manage large tables by distributing them across filegroups. - Partitioning Key: Determines how the data is divided.
Sparse Files	Files that only allocate disk space as data is written, reducing the initial size of the file.	<ul style="list-style-type: none"> - Purpose: Saves space when the actual data stored is sparse. - Use Case: Typically used for columns that have a lot of NULL values.
Instant File Initialization	A feature that allows SQL Server to skip zeroing out space when creating or growing data files, speeding up operations.	<ul style="list-style-type: none"> - Benefits: Speeds up file creation and growth by not initializing the file with zeros. - Requirement: Requires specific Windows permissions (Perform volume maintenance tasks).
File Growth	The automatic expansion of a data file when it reaches its allocated size.	<ul style="list-style-type: none"> - Settings: Can be configured to grow by a fixed size or by a percentage. - Impact: Uncontrolled growth can lead to fragmentation or excessive VLFs in the transaction log.
Shrink Operations	The process of reducing the size of a data file by removing unused space.	<ul style="list-style-type: none"> - Consideration: Frequent shrinking can cause fragmentation and impact performance. - Best Practice: Should be used cautiously and typically avoided unless necessary.
File Stream	A feature that allows large binary data (e.g., documents, images) to be stored directly on the file system instead of in the database.	<ul style="list-style-type: none"> - Integration: Managed through T-SQL but stored outside the database files. - Use Case: Ideal for storing large unstructured data with transactional consistency.

Max Size	The maximum size a data file can reach, configurable during file creation or through later modifications.	<ul style="list-style-type: none"> - Settings: Can be set to a fixed maximum or to unlimited (auto-grow until disk is full). - Consideration: Proper planning is required to avoid running out of space unexpectedly.
TempDB Architecture	The system database used for temporary objects, intermediate query results, and version stores.	<ul style="list-style-type: none"> - Files: Usually has multiple data files for better performance in multi-core systems. - Rebuilt on Startup: TempDB is recreated every time SQL Server restarts, clearing all temporary data.
Buffer Pool	The area of memory where SQL Server caches pages read from the database files.	<ul style="list-style-type: none"> - Purpose: Reduces the need for disk I/O by keeping frequently accessed data in memory. - Lazy Writer: Process that periodically writes dirty pages back to disk.

This table provides a comprehensive overview of SQL Server Data File Architecture, covering the structure, management, and operational aspects of data storage within SQL Server.

<https://www.sqlbadchamps.com>